

# Relational Databases II

**License:**

Health Information Technology by Hye-Chung Kum is licensed under a [Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-nc-sa/4.0/)

**Course URL:**

<http://pinformatics.org/phpm631>



POPULATION  
INFORMATICS  
RESEARCH GROUP



# Outline

- What is a database?
- What is a database management system?
- An Introduction to SQL
  - How to retrieve data from a database I
    - Basic SQL Queries
  - How to retrieve data from a database II
    - Calculations using multiple columns or rows
  - How to retrieve data from a database III
    - Combining tables
  - How to create a database (optional)

# Table Operations

- Aggregate columns: col1 op col2 **AS** col3

col1	col2	→	col1	col2	col3
a	d		a	d	a+d
b	e		b	e	b+e
c	f		c	f	c+f

- Aggregate rows: **Group BY**

A	→	D
B		
C		

Where  $D = \text{function}(A, B, C)$

Examples of function are

Sum(A,B,C) Avg(A,B,C) Max(A,B,C) Min(A,B,C) Count(A,B,C)

# Compute Columns

- Find discount amount
  - SELECT patientID, (billed-covered) AS discount
  - FROM payments;
- Nice names for output columns
  - Name following computed column (e.g., discount) will be used to name output column
- String vars: concatenate ||
  - fname || " " || lname AS name
- Find total paid amount (payments Table)
  - Total = copay+pat\_pd+insur\_pd

```
SELECT calculate AS NewColumnName  
FROM Table ;
```

# Grouping and Aggregates I

- Can make calculations on groups of rows
  - sum, avg, max, min, count
- Each different value for the GROUP BY fields defines a new group
  - One row of output is produced for each group
  - Several rows of input table may belong to same group. They are aggregated using aggregation operator.

```
SELECT f(Column2) AS ColumnName  
FROM Table  
GROUP BY Column1;  
--f(x) : sum, avg, max, min, count ;
```

# Grouping and Aggregates II

- Can make calculations on groups of rows
  - sum, avg, max, min, count
- How many visits did each patient have?
  - SELECT patientID, count(visitID) AS nvisits
  - FROM visits
  - GROUP BY patientID;
- TRY: What is total billed by patient?
  - Payments table

```
SELECT f(Column2) AS ColumnName
FROM Table
GROUP BY Column1;
--f(x) : sum, avg, max, min, count ;
```

# Take Away II

## SQL – Structured Query Language

- Every statement yields a table of values as output
  - Sometimes there's only one row in the table!

Keyword	parameters
<b>select</b>	columns and/or expressions ( <b>AS</b> )
<b>from</b>	Tables
<b>where</b>	conditions on the rows
<b>group by</b>	group rows together
<b>order by</b>	order the rows
<b>;</b>	

# Outline

- What is a database?
- What is a database management system?
- An Introduction to SQL
  - How to retrieve data from a database I
    - Basic SQL Queries
  - How to retrieve data from a database II
    - Calculations using multiple columns or rows
  - How to retrieve data from a database III
    - Combining tables: Joins
  - How to create a database (optional)



# Joins

- Combine rows from one table with rows from another
- Usually join on some common column
  - Don't combine rows unless their value in the common column is the same
  - WHERE clause says the common column must be same in each table
- Give discount amount by last name of patient
  - SELECT patientID, patients.lname, (billed-covered) AS discount
  - FROM payments, patients
  - Where payments.patientID=patients.patientID;
- Give primary diagnosis (diag1) description for each patient visit
  - Tables Visits & lu\_diag

```
SELECT Table1.Column1, Table2.Column2
FROM Table1, Table2
WHERE Table1.Column=Table2.Column;
```

# Practice Problems

- Give primary diagnosis (diag1) description for each patient visit
  - Tables Visits & lu\_diag
- P1: for only patients that saw Dr. Gaines
- P2: for only patients seen in dec 2012
- P3: for only patients seen in dec 2012 by Dr. Gaines
- P4: for only patients that saw Dr. Gaines or Dr. Fry

# Different Syntax: Joins

```
SELECT Table1.Column1, Table2.Column2  
FROM Table1, Table2  
WHERE Table1.Column=Table2.Column;
```

-- Does same thing;

```
SELECT Table1.Column1, Table2.Column2  
FROM Table1  
JOIN Table2  
ON Table1.Column=Table2.Column;
```



# Different SQL JOINS

- INNER JOIN: Returns all rows when there is at least one match in **BOTH tables**
- LEFT JOIN: Return **all rows from the left table**, and the matched rows from the right table
- RIGHT JOIN: Return **all rows from the right table**, and the matched rows from the left table
- FULL JOIN: Return all rows when there is a match in **ONE of the tables**



# Views: Permanent Queries

- Looks and feels like a table
- Saved queries
- Virtual table: not a real table in the DB
- Can treat it like a real table, as if it exists



POPULATION  
INFORMATICS  
RESEARCH GROUP



# Create View

```
CREATE VIEW panel as  
SELECT
```

```
    providers.fname AS dr_first,  
    providers.lname AS dr_last,  
    patients.fname,  
    patients.lname,  
    patientID
```

```
FROM providers, patients
```

```
WHERE
```

```
    providers.providerID=patients.primary_dr
```

```
ORDER BY
```

```
    providers.providerID;
```



# Practice Problems

- How many patients does each doctor have?
  - Hint: use the view you just created
- How did you identify the doctor?
  - Last name? first name?
  - What if there are two doctors with the same name?
  - What is the BEST way to refer to entities in a RDB?



# Create (Optional)

- Primary Key
- Data Type
  - Text, Integer, Real
  - Data types might have different names in different database. And even if the name is the same, the size and other details may be different! **Always check the documentation!**
  - [http://www.w3schools.com/sql/sql\\_datatypes\\_general.asp](http://www.w3schools.com/sql/sql_datatypes_general.asp)
- Example: kumdb.sql





# Advanced conditionals (Optional)

## Like & Wildcard

Wildcard	Description
%	A substitute for zero or more characters
_	A substitute for a single character
[charlist]	Sets and ranges of characters to match
[^charlist] or [!charlist]	Matches only a character NOT specified within the brackets

```
SELECT fname, lname  
FROM patients  
WHERE lname LIKE 's%';
```



# Indexing

MEMORY

00	901
01	398
02	901
03	399
04	598
05	199
06	902
07	000
...	
98	
99	

BLACKBOARD


w  
b

- Can have many per table
- Mapping to locations in storage, for quick lookup
  - Example: index patientID
    - PID=1 in memory 04
    - PID=2 in memory 08 ...etc...
- Pros: Faster to find and retrieve data
- Cons: Slow to enter and save data



# ACID: Atomic, Consistent, Isolated & Durable

- four most desirable basic characteristics of a transaction system
- **Atomic** transactions are such that the transaction is either entirely completed or makes no change to the database; even if an error or a hardware fault occurs mid-transaction the database will not be left with a half-completed transaction.
- **Consistent** transactions ensure that the database is left in a consistent state after the transaction is complete, meaning that any integrity constraints (unique keys, foreign keys, CHECK constraints etc.) must be satisfied or the transaction will be rejected.
- **Isolated** transactions are invisible to other users of the database while they are being processed
- **Durable** transactions guarantee that they will not be rolled back after the caller has committed them
- Early RDBMSes couldn't always guarantee all four of these requirements with their transactions, but modern counterparts can usually provide ACID transactions even in the event of power or hardware failure.



# Take Away I

## Advanced SQL queries

- Why write SQL queries?
  - To answer real world questions from the database
- Calculating new columns from other columns
- Combining multiple rows
  - Grouping and aggregating
- Views: Permanent Queries
  - **CREATE VIEW**

# Take Away II

## Different SQL JOINS

- INNER JOIN: Returns all rows when there is at least one match in **BOTH tables**
- LEFT JOIN: Return **all rows from the left table**, and the matched rows from the right table
- RIGHT JOIN: Return **all rows from the right table**, and the matched rows from the left table
- FULL JOIN: Return all rows when there is a match in **ONE of the tables**



# Take Away III

MEMORY

00	901
01	398
02	901
03	399
04	598
05	199
06	902
07	000
...	
98	
99	

BLACKBOARD

w  
b

- Indexing
  - Can have many per table
  - Mapping to locations in storage, for quick lookup
    - Example: index patientID
      - PID=1 in memory 04
      - PID=2 in memory 08 ...etc...
  - Pros: Faster to find and retrieve data
  - Cons: Slow to enter and save data
- four most desirable characteristics of a transaction system
  - ACID: Atomic, Consistent, Isolated & Durable

